

90 Hypotheses, Zero Edges

I built a pre-registered research pipeline to find tradeable edges for a £1,500 account. It correctly found nothing. This is the short version of the audit trail.

Over several months I tested roughly 100 ideas across prediction markets, corporate actions, index flows, insider filings and ETF cost structures. None of them pays at that scale. That is the result. Everything below is why I believe it, which is the only part worth reading.

Why a negative result is worth anything

Most retail backtests find edges, and almost all of those edges are artifacts. The question is never "did you find something" but "would you have known if you hadn't." Five things made the answer trustworthy.

1. **Pre-registration with a stated direction.** Each hypothesis was registered with its predicted sign before data was pulled, hash-chained so it cannot be revised afterwards. A sentiment signal registered as "attention predicts returns" came back at $t = 0.04$. Without pre-registration, the honest move — reporting nothing — is also the least appealing one.
2. **Correcting for how many tests I ran.** With $m = 38$ hypotheses at FDR $q = 0.05$, the Benjamini-Hochberg bar is $p = 0.00685$, not 0.05 . Testing 155 variants at $p < 0.05$, as one popular method does, yields ~ 7.8 false positives by chance alone — so "5 survivors from 155" is below what pure noise produces.
3. **Overlap-aware inference.** With 63–252 day holding periods entered continuously, dozens of positions are open at once and driven by the same market. Treating them as independent overstates the evidence by $17\text{--}27\times$. Collapsing to calendar-time monthly returns moved one hypothesis from $p = 0.049$ to $p = 0.87$.
4. **Factor attribution before celebrating.** A buyback-announcement strategy showed 14.5% raw. Against a six-factor model: alpha 4.6%, beta 1.00, $R^2 = 0.96$ — almost entirely market exposure. Separately, the three largest "alphas" in the batch had $R^2 = 0.00$: failed regressions, not discoveries. Ranking by alpha without checking R^2 puts the most broken results on top.
5. **Bounds instead of assuming the missing data away.** Where $\sim 46\%$ of events were unobservable, a Manski bound gives what the result could be under any assumption about the missing half, rather than silently conditioning on survivors.

The phantom £74,000

My tender scanner reported fourteen profitable candidates worth \$79,978. Validated against itself, **98.6% of that profit was artifacts**. \$74,417 of it came from one position: a tender price of \$850 compared against a \$94.53 market price that had been split-adjusted for a 10-for-1 split occurring after the tender — an implied spread of roughly 800%, which is not a number any real tender produces. Another \$4,448 was a Canadian-dollar tender priced against the US listing.

Both are the same error in different clothes: a price compared against a price that was never the same thing. Neither was subtle once seen; neither was visible in aggregate. Real candidates: eleven,

worth \$1,113.

Twenty-one defects, none caught by the tests

I kept a defect catalogue — one row per defect, recording what I had believed, what the effect was, and what caught it. Twenty-one entries over five weeks, and the interesting column is what is absent from it: **not one defect was caught by the test suite**. They were caught by hand-reading real filings, by spot-checking companies known to be alive, by opening a query form and getting a subscription error, by noticing one ticker flip between runs, by asking why a risk-free strategy had six stop-loss exits.

One entry is worse than the rest. The nightly digest reported "420 tests OK". Full discovery was 833 tests with 17 failures — the green number was measured on a six-file subset that excluded almost everything that session had changed. The suite was not merely failing to catch defects; it was reporting green over red, on precisely the files under active modification.

The turn

One deal was worth \$176.32 against \$621.56 earned across every tradeable deal in three years — 28% of the strategy, lost to a missing word and an italic tag in a filing I had successfully downloaded, parsed and discarded. I spent two fixes recovering it.

In the same audit sat an unrelated finding: the price parser was reading Dutch auction ranges backwards. That deal was a Dutch auction, range \$20.75 to \$23.75, market price \$21.96. Scored against \$23.75 it is worth +\$176. Scored against \$20.75, which is what the method always said it used, the same deal is worth −\$120. **The prize behind the bug was a loss**, and the figure that made both fixes feel urgent was itself produced by a third defect the same morning was in the middle of finding.

What this is actually about

Dropping filters are the dangerous class. A bug that invents a deal is caught the moment you read the filing. A bug that discards one is indistinguishable from an honest empty market, and it is self-concealing: the quieter your scanner gets, the more reasonable it looks.

Fixtures only verify what you already understood. A fixture is a snapshot of your imagination on the day you wrote it. My tests passed throughout — they were testing that the scanner still did what I had wrongly assumed it should. Only contact with unfamiliar real documents moves that boundary.

What I would tell someone starting

- Ask whether the thing is even observable before writing the scanner. It is the cheapest filter available and it dominates every statistical control that comes after.
- Compute your multiple-testing bar first. If you cannot state m , you are not testing hypotheses, you are shopping.

- Check every aggregate one level finer before believing it — filings against companies, rows against episodes. I made this error seven times in one day, and every time the wrong number was the flattering one.
- A statistic computed at the wrong grain arrives with the authority of real data. That makes it more dangerous than a guess, not less.
- Verify the execution path before you verify the edge. Ten minutes with broker support would have reordered eight months of work.

The pipeline did its job. It killed almost everything, honestly, including several results I had already written up as wins. The methodology is the deliverable; the P&L is the evidence it was applied.

Nothing described here traded real money. All results are paper or backtest. This is a research log, not advice. — Charlie Wuytack, August 2026